

CFEngine Zero to Hero Primer

Nick Anderson

nick@cmdln.org

Introduction

CFEngine contains a powerful language for controlling all aspects of a system. CFEngine runs primarily on UNIX and UNIX like operating systems, but can also run on Windows.

CFEngine is very extensive and powerful. In this presentation you will learn only a subset of what CFEngine can do. A mere tip of the iceberg, but this will represent the bulk of what you do with CFEngine. In other words, you'll learn the 20% of CFEngine that will do 80% of the work.

Want to know more?

- [Vertical Sysadmin's CFEngine courses](#)
- [CFEngine Reference Manual](#)
- [CFEngine Training & Professional Services](#)
- [Beyond Automation with CFEngine3](#) (Video Training)

Fork Me on Github!

You can get a copy of this presentation any time on Github.

<http://github.com/nickanderson/CFEngine-zero-to-hero-primer>

CFEngine Components

These are the major components of CFEngine that you will encounter on a day to day basis.

- `cf-agent`
- `cf-monitord`
- `cf-execd`
- `cf-serverd`

cf-agent

`cf-agent` is the command you will use most often. It is used to apply policy to your system. If you are running any CFEngine command from the command line, there's a greater than 99% chance that this is it.

cf-monitord

cf-monitord monitors various statistics about the running system. This information is made available in the form of **classes** and **variables**.

You'll almost never use cf-monitord directly. However the data provided by cf-monitord is available to cf-agent.

cf-execd

cf-execd is a periodic task scheduler. You can think of it like cron on steroids. By default CFEngine runs and enforces policies every *five minutes*. cf-execd is responsible for making that happen.

cf-serverd

cf-serverd runs on the CFEngine server, as well as all clients.

- On servers it is responsible for serving files to clients.
- On clients it accepts cf-runagent requests

cf-runagent allows you to request ad-hoc policy runs. I rarely use it.

How vs. What



Imperative vs. Declarative

It is very likely that you have only ever used **imperative** languages. Examples of imperative languages include C, Perl, Ruby, Python, shell scripting, etc. Name a language. It's probably imperative.

CFEngine is a **declarative** language. The CFEngine language is merely a *description* of the final state. CFEngine uses **convergence** to arrive at the described state.

Imperative is Sequential

Imperative languages execute step by step in **sequence**.

Because it is sequential must go in order from start to finish. If a sequential program is interrupted in the middle the state is *inconsistent*. Neither at starting point nor at the finish. Executing the program again will repeat tasks that have already been done, possibly causing damage in doing so.

E.g., a script that appends a line to a file then restarts a daemon. If the line is appended but the daemon not restarted, running the script again will result in that line in the file twice. This will likely cause the daemon to fail when restarted.

Imperative starts at known state A and transforms to known state B.

Declarative is Descriptive

It is not a list of steps to arrive achieve an outcome but a **description** of the desired state. Because of this any deviation from the desired state can be detected and corrected.

In other words, a declarative system can begin in *any* state, not simply a known state, and transform into the desired state.

Declarative states a list of things which must be true. It does not state how to make them true.

When a system has reached the desired state it is said to have reached **convergence**.

Promise Theory

Promise theory is the **fundamental underlying philosophy** that drives CFEngine.

It is a model of voluntary cooperation between individual, autonomous actors or agents who publish their intentions to one another in the form of promises.

What makes promises?

A file (e.g., `/etc/apache2/httpd.conf`) can make promises about its own contents, attributes, etc. But it does not make any promises about a process.

A process (e.g., `httpd`) can make a promise that it will be running. But it does not make any promises about its configuration.

The configuration file and the process are *autonomous*. Each makes promises about itself which cooperates toward an end.

Going Deeper


- Thinking in Promises
- In Search of Certainty
- Promise Theory: Principals and Applications

O'REILLY®

"Truly, a blueprint for the systems of tomorrow."
—Mike Dvorkin, Distinguished Engineer at Cisco

THINKING IN PROMISES

DESIGNING SYSTEMS FOR COOPERATION



MARK BURGESS

O'REILLY®


"An instant classic in computer science..." —Glenn O'Donnel

2nd Edition
Revised & Updated

In Search of Certainty

The Science of
Our Information
Infrastructure

Mark Burgess



Promises

Anatomy of a Promise

```
type:  
  context::  
    "promiser" -> "promisee"  
    attribute1 => "value",  
    attribute2 => value2;
```

- **type** is the kind of promise being made (e.g., files, commands, etc.).
- **context** is optional and defaults to `any ::`. Promises with a context will only apply if the given context is true.
- **promiser** is what is making the promise. (e.g., a file or a process).
- **promisee** is an optional recipient or beneficiary of the promise.

Promise Attributes

```
type:  
  context::  
    "promiser" -> "promisee"  
    attribute1 => "value",  
    attribute2 => body;
```

Each promise can have one or more attributes that describe the parameters of the promise. The available attributes will vary depending on the **promise type**.

The value can be either a text string (which must be quoted) or another object (which must not be quoted). All of the attributes together are called the **body** of the promise (as in "the body of an e-mail", or "the body of a contract").

Attributes are separated by **commas**. Each promise ends with a **semi-colon**.

Example Promise

```
files:  
  linux::  
    "/tmp/hello/world" -> { "Student" }  
    create => "true";
```

- This is a promise of **type** files.
- This promise has a **class context** of `linux` (it will only apply if running a Linux kernel).
- The **promiser** is the POSIX path `/tmp/hello/world`.
- This promise has only one **attribute**, specifying that the file should be created if it does not exist.
- The **promisee** is *you!*
- To create a directory instead, use a `files:` promise and append a `.` to the directory name (e.g., `/tmp/hello/.`)

Bundles

A **bundle** is a collection of **promises**. It is a logical grouping of any number of promises, usually for a common purpose. E.g., a bundle to configure everything necessary for Apache to function properly.

For example, a bundle to configure Apache might:

- install the apache2 package
- edit the configuration file
- copy the web server content
- configure filesystem permissions
- ensure the httpd process is running
- restart the httpd process when necessary

Anatomy of a Bundle

```
bundle type name
{
    type:
    context::
        "promiser" -> "promisee"
        attribute1 => "value",
        attribute2 => value;

    type:
    context::
        "promiser" -> "promisee"
        attribute1 => "value",
        attribute2 => value;
}
```

Bundles apply to the binary that executes them. E.g., agent bundles apply to cf-agent while server bundles apply to cf-serverd.

Bundles of type common apply to any CFEngine binary.

For now you will only create agent or common bundles.

Bodies

I stated before that the attributes of a promise, collectively, are called the body. Depending on the specific attribute the value of an attribute can be an **external body**.

A **body** is a collection of *attributes*. These are attributes that supplement the promise.

Anatomy of a Body

```
body type name {  
  attribute1 => "value";  
  attribute2 => "value";  
}
```

The difference between a *bundle* and a *body* is that a bundle contains *promises* while a *body* contains only *attributes*.

Take a moment to let this sink in.

- A **bundle** is a collection of *promises*.
- A **body** is a collection of *attributes* that are applied to a promise.

The distinction is subtle, especially at first and many people are tripped up by this.

In a body each attribute ends with a **semi-colon**.

Abstraction and Re-usability

Bundles and bodies can be parameterized for abstraction and re-usability. In other words you can define one and call it even passing in parameters which will implicitly become variables.

Example

```
body type name (my_param) {  
  attribute1 => "$(my_param)";  
}
```

The parameter `my_param` is accessed as a variable by `$(my_param)`.

The Masterfiles Policy Framework

The **Masterfiles Policy Framework** is the default policy that ships with CFEngine. The standard library is included.

- Masterfiles Policy Framework

CFEngine Standard Library

The **CFEngine Standard Library** comes bundled with CFEngine in the `masterfiles/lib/` directory.

The standard library contains ready to use bundles and bodies that you can include in your promises and is growing with every version of CFEngine. Get to know the standard library well, it will save you much time.

- [Standard Library Reference](#)

Putting it All together

These are the building blocks. You now know what they all are.

Examples

Now we will go through some examples.

I encourage you to try executing the examples as we go along.

- <https://github.com/nickanderson/CFEngine-zero-to-hero-primer/tree/master/examples>

To execute a policy run the following command:

```
$ cf-agent --file ./test.cf --bundle bundlename
```

Note: Make sure you use the correct file and bundle name! For any examples using a bundle named main you can skip specifying the bundle.

Running commands

commands_echo_hello_world.cf

```
bundle agent main
{
  commands:
    "/bin/echo Hello World!";
}
```

Set File Permissions

set_file_permissions.cf

```
bundle agent example {  
  files:  
    "/etc/shadow"    perms => perms_for_shadow_files;  
    "/etc/gshadow"   perms => perms_for_shadow_files;  
}  
  
body perms perms_for_shadow_files {  
  owners => { "root" };  
  groups => { "shadow" };  
  mode   => "0640";  
}
```

- This is an **agent** bundle (meaning that it is processed by cf-agent).
- Its purpose is to set the permissions on /etc/shadow and /etc/gshadow.
- It uses an external body named perms_for_shadow_files.
- The body only needs to be defined once and can be reused for any number of promises.

Note: The values for owners and groups is enclosed in curly braces. This is because these attributes take a list of strings (aka, an slist).

Copy an Entire File

```
bundle agent example {  
  files:  
    "/etc/motd"    copy_from => cp("/repo/motd");  
}  
  
body copy_from cp (from) {  
  servers    => { "$(sys.policy_hub)" };  
  source     => "$(from)";  
  compare    => "digest";  
}
```

- The purpose of this bundle is to copy /etc/motd from the CFEngine server
- `$(sys.policy_hub)` is an automatic variable which contains the CFEngine server's address.
- The path `/repo/motd` is on the *server's* filesystem.
- The `compare` type tells CFEngine how to know when the file needs updating.

Edit a File

sshd_permit_root_login_no.cf

```
bundle agent main {
  files:
    "/etc/ssh/sshd_config"    edit_line => deny_root_ssh;
}

bundle edit_line deny_root_ssh {
  delete_lines:
    "^PermitRootLogin.*"
  insert_lines:
    "PermitRootLogin no"
}
```

- This will delete any line matching the regular expression `^PermitRootLogin.*`.
- This also inserts the line `PermitRootLogin no` **at the end of the file**.
- CFEngine is smart enough to know not to edit the file if the end result is already *converged*.
- This is an overly simplistic example. When editing configuration files you probably want to copy the whole file or use `set_config_values()` from the standard library.

Classification and Classes

A **class** is like a tag (like tagging a photo). Classes are used to give a promise **context**. There are two types of classes.

1. **Built in classes.** These so called **hard classes** are classes that CFEngine will create automatically. Hard classes are determined based on the system attributes. For example a server running Linux will have the class `linux`.
2. **User defined classes.** These so called **soft classes** are classes that are defined by you. You can create them based on the outcome of a promise, based on the existence of other classes, or for no reason.

My classes

Here is a list of hard classes defined on an actual system running CFEngine.

```
cf-promises --show-classes
```

Class name

127_0_0_1

172_17_0_1

192_168_42_189

4_cpus

64_bit

Afternoon

Day7

GMT_Day7

GMT_Evening

GMT_Hr21

GMT_Hr21_Q3

GMT_Lcycle_0

GMT_May

GMT_Min40_45

GMT_Min43

GMT_Q3

GMT_Saturday

GMT_Yr2016

Hr16

Hr16_Q3

Lcycle_0

May

Min40_45

Min43

Q3

Saturday

Yr2016

any

cfengine

cfengine_3

cfengine_3_8

cfengine_3_8_1

common

compiled_on_linux_gnu

Meta tags

inventory,attribute_name=None,source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

source=agent,derived-from=sys.cpus,hardclass

source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

time_based,source=agent,hardclass

source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

inventory,attribute_name=None,source=agent,hardclass

cfe_internal,source=agent,hardclass

source=agent,hardclass

Control Promise Selection

```
bundle agent apache_config {  
  files:  
  
    debian::  
      "/etc/apache2/apache2.conf"  
      copy_from => remote_cp("/cfengine/repo/debian/apache2.conf", "${sys.policy_hub}");  
    redhat::  
      "/etc/httpd/conf/httpd.conf"  
      copy_from => remote_cp("/cfengine/repo/redhat/httpd.conf", "${sys.policy_hub}");  
    solaris::  
      "/etc/apache2/2.2/httpd.conf"  
      copy_from => remote_cp("/cfengine/repo/solaris/httpd.conf", "${sys.policy_hub}");  
}
```

This set of promises will copy the appropriate apache config file depending on the type of server. Notice that each file promise is prefixed by a **class**. The promise will be skipped unless that class is defined on the system.

Thus, only Debian systems will run the `debian::` context promise, only Red Hat will run `redhat::` and only Solaris will run `solaris::`.

Promise Type and Class Context Can be Implicit

The *promise type* and *class context* don't need to be listed for every promise. Think of each like a heading in an outline. Everything that follows is still under the same heading until a new heading is declared. If a new promise type is declared the class context is reset as well.

implicit_class_context.cf

```
bundle agent example {
  files:
    solaris::
      "/tmp/hello/world"
      create => "true";
      "/tmp/foo/bar"
      create => "true";
    linux::
      "/dev/shm/hello_world"
      create => "true";
  commands:
    "echo Hello World";
}
```

The first three promises are of type `files`. The first two will only execute on `solaris` while the third will only execute on `linux`. The last promise has a new promise type, of `commands`, and will always execute.

A Note About Classes and Distributions Based on Other Distributions

I said that only Debian systems will run `debian::` and only Red Hat will run `redhat::`. This isn't exactly true.

- Ubuntu is based on Debian, and so will have both `ubuntu` and `debian` defined as hard classes.
- Likewise, CentOS is based on Red Hat and so will have both `centos` and `redhat` defined as hard classes.

This goes for any distro that is based on another distro. The "parent" classes will be also defined.

Use Classes to Control Flow

```
bundle agent apache_config {
  files:

    "/etc/apache2/apache2.conf"
    copy_from => remote_cp("/cfengine/repo/debian/apache2.conf", "${sys.policy_hub}")
    classes => if_repaired("RestartApache");

  commands:

    RestartApache::
      "/usr/sbin/apache2ctl graceful";
}
```

This set of promises will first copy the Apache configuration file. Once the Apache configuration file is updated, Apache must be restarted. In order to make sure that Apache gets restarted when necessary a class will be defined when the configuration file is updated.

When CFEngine reaches the commands section, if the RestartApache class is defined (which only happens if the config file is updated) then Apache will be restarted.

Use Classes to Control Flow

```
bundle agent apache_config {
  files:

    "/etc/apache2/apache2.conf"
    copy_from => remote_cp("/cfengine/repo/debian/apache2.conf", "${sys.policy_hub}"),
    classes => if_repaired("RestartApache");

  commands:

    RestartApache::
      "/usr/sbin/apache2ctl graceful";
}
```

So, the workflow then is:

1. Perform promise 1
2. Define a class if repaired
3. Perform promise 2 if the class has been set

I use this **ALL. THE. TIME**. If this class is to teach you 20% that accomplishes 80%, **this slide** is the 5% that accomplishes 95%.

Class Expressions

```
commands:  
  RestartApache.debian::  
    "/usr/sbin/apache2ctl graceful";  
  RestartApache.redhat::  
    "/usr/sbin/apachectl graceful";
```

This example is similar to the last one, except that Debian and Redhat each have different commands used to restart Apache. Therefore, we use an expression to define our class context. The expression `RestartApache.debian` means "RestartApache **and** debian".

Class Expressions

commands:

```
RestartApache.debian::  
    "/usr/sbin/apache2ctl graceful";  
RestartApache.redhat::  
    "/usr/sbin/apachectl graceful";
```

Operator	Meaning	Example
. and &	boolean and	debian.Tuesday::
and	boolean or	Tuesday Wednesday::
!	boolean not	!Monday::
()	Explicit grouping	(debian redhat).!ubuntu.!centos::

Managing Processes

Keep Services Running: Using Processes

```
bundle agent apache {  
    processes:  
        "apache2"  
        restart_class => "StartApache";  
  
    commands:  
        StartApache::  
            "/etc/init.d/apache2 start";  
}
```

This policy uses a processes promise to check the process table (with ps) for the regular expression `. *apache2 . *`. If it is not found then the class `StartApache` will get defined.

When CFEngine executes commands promises Apache will be started.

Ensuring Processes are Not Running: Using Processes and Commands

process_stop_bluetoothd.cf

```
bundle agent stop_bluetooth {  
  processes:  
    "bluetoothd"  
    process_stop => "/etc/init.d/bluetooth stop";  
}
```

This policy uses a `processes` promise to check the process table (with `ps`) for the regular expression `.*bluetoothd.*`. If it is found the `process_stop` command is executed.

Ensuring Processes are Not Running: Using Processes and Signals

process_signals_bluetoothd.cf

```
bundle agent stop_bluetooth {  
  processes:  
    "bluetoothd"  
    signals => { "term", "kill" };  
}
```

This policy uses a `processes` promise to check the process table (with `ps`) for the regular expression `.*bluetoothd.*`. Any matching process is sent the `term` signal, then sent the `kill` signal.

Note: The promise `bluetoothd` becomes the **regular expression**, `.*bluetoothd.*` that is matched against the output of `ps`. This means that it can match **anywhere** on the line (in versions prior to 3.9), not just the process name field. **Caveat emptor!**

Keep Services Running: Using Services

```
bundle agent apache {  
  services:  
    "www"  
    service_policy => "start";  
}
```

This uses the `services` promise type to ensure that Apache is always running.

The `standard_services` bundle implementation currently covers `systemd`, `chkconfig`, the `service` command, `svcadm` and `systemV` init scripts. Proper functionality relies on each installed service correctly implementing a service check as appropriate for the init system in use.

Ensuring Processes are Not Running: Using Services

services_bluetoothd_stop.cf

```
bundle agent stop_bluetoothd {  
  services:  
    "bluetoothd"  
    service_policy => "stop";  
}
```

This policy uses a `services` promise type to ensure that Bluetooth services are not running. Again, this only works for services that are defined under `standard_services` in the standard library and requires `cfengine` 3.4.0 or higher.

The same restrictions about distros apply to stopping services promises.

Managing Packages

Legacy Implementation

```
bundle agent install {  
  packages:  
    "zsh"  
    package_policy => "addupdate",  
    package_method => apt,  
    package_select => ">=",  
    package_version => "4.3.10-14";  
}
```

- The `package_policy` of `add update` will install or upgrade. Using `add` will only install, never upgrade, `upgrade` will upgrade only and `delete` will uninstall.
- The `package_method` of `apt` is in the standard library, look there for other package methods (e.g., `rpm`, `ips`, etc.).
- The `package_select` of `>=` means the installed version must be equal to or newer than the specified version or it will be replaced. Using `<=` would downgrade, if the `package_method` supports downgrading and `==` will require the exact version.

New Implementation

```
bundle agent install {  
  packages:  
    "zsh"  
    policy => "present",  
    package_module => yum,  
    version => "latest";  
}
```

- The `policy` of `present` will make sure the package is installed on the system, while a `policy` of `absent` will ensure a package is not installed.
- The `package_module` of `yum` is included in the Masterfiles Policy Framework.
- The `version` of `latest` means the installed version should be the latest available. Alternatively you can provide an explicit version.

Package Managers

- package_methods

`pip(flags), npm(dir), npm_g, brew(user), apt, apt_get, apt_get_permissive, apt_get_release(release), dpkg_version(repo), rpm_version(repo), windows_feature, msi_implicit(repo), msi_explicit(repo), yum, yum_rpm, yum_rpm_permissive, yum_rpm_enable_repo(repo_id), yum_group, rpm_filebased(path), ips, smartos, smartos_pkg_add(repo), opencsw, solaris(pkgname, spoolfile, adminfile), solaris_install(adminfile), freebsd, freebsd_portmaster, alpinelinux, emerge, pacman, zypper, generic`

- package_modules

`yum, apt_get, freebsd_ports, nimclient, pkg, pkgsrc`

Managing Files

Templating a file

template.mustache

```
Hello from {{{vars.sys.fqhost}}}!
```

```
{{#classes.linux}}I am a Linux Box!{{/classes.linux}}  
{{^classes.windows}}I am NOT a Windows Box{{/classes.windows}}
```

template_file.cf

```
bundle agent main{  
  files:  
    "/tmp/example"  
    create => "true",  
    edit_template => "${this.promise_dirname}/template.mustache",  
    template_method => "mustache";  
}
```


Deleting a file

```
bundle agent tidy {  
  files:  
    "/var/log/*"  
    file_select => days_old("7"),  
    delete => tidy;  
}
```

This policy will delete any files in `/var/log/` older than 7 days. The `days_old()` and `tidy` bodies are included in the standard library,

To delete a file indiscriminately, omit the `file_select`.

Look up `file_select` and `tidy` in the [reference-manual](#) to find more ways to use this.

Setting Up a Client/Server Environment

Before starting you need to have cfengine installed on the server and the client and the server FQDN must be set properly in DNS (or use the IP addresses). This is ideally handled by your provisioning process. Along with automating server function you should also be automating your provisioning process.

Some ways of automating provisioning are `kickstart`, `preseed`, `fai`, `cobbler`, `disk imaging`, `instance cloning`, etc, etc. This of course is not a complete list.

Bootstrapping the Server and Client

Server Side

Edit `/var/cfengine/masterfiles/def.cf` to set the `acl` list for the IP addresses of your network, then run:

```
cf-agent --bootstrap $(hostname --fqdn)
cf-agent -KI
```

Client Side

Simply run:

```
cf-agent --bootstrap server.fqdn.example.com
```

You can use the server's IP address instead of the DNS name.

Managing and Distributing Policies

The policy files are in `/var/cfengine/masterfiles` on the server (also known as the `policy_hub`) and are copied to `/var/cfengine/inputs`. All clients then copy `/var/cfengine/inputs` from the server.

<div style=text-align:center"></div>

Now edit the policy in `/var/cfengine/masterfiles` on the server and watch for the changes to happen on the client.

As you write new policies, each bundle needs to be listed in the `bundlesequence` and each file needs to be listed in `inputs`. Both of these are under `body common control` inside of `promises.cf`.

Bundles are executed in the order they are listed in the `bundlesequence`, but `inputs` can be listed in any order.

Reporting on Promise Outcomes

CFEngine logs to `/var/cfengine/promise_summary.log`. Here's an example log message:

```
1463018982,1463018990: Outcome of version CFEngine Promises.cf 3.7.0 (agent-0):\nPromises observed - Total promise compliance: 93% kept, 3% repaired,\ 4% not kept (out of 148 events).\nUser promise compliance: 93% kept, 2% repaired, 5% not kept (out of 130 events).\nCFEngine system compliance: 94% kept, 6% repaired, 0% not kept (out of 18 events).
```

Note: The timestamp is a Unix epoch.

CFEngine will also send an email to the configured address in `body executor control=` any time there is output from an agent run that differed from the previous run.

And finally you can use the `-I` flag to have CFEngine **inform** you of repairs. (Shown here along with the `-K` flag which ignores any lock timers).

```
cf-agent -KI
```

Enterprise Reporting

CFEngine

4 hostsHealth OKHA status N/AHello, admin

Dashboard

Hosts

Reports

Monitoring

Design Center beta

Results

> Filters

Promise outcome of last run

Showing 50 of 510 results on 11 pages

12345»Last50

Host name	Bundle name	Promise type	Promiser	Promise handle	Promise outcome	Log messages	Change time
hub	cfe_internal_bins	files	/usr/local/sbin/cf-monitord.cfsaved	null	KEPT		2016-05-17 13:13:33+00
hub	cfe_internal_bins	files	/usr/local/sbin/cf-runagent.cfsaved	null	KEPT		2016-05-17 13:13:33+00
host001	autorun	methods	autorun	null	KEPT		2016-05-17 14:41:55+00
host001	broken_promises	meta	tags	null	KEPT		2016-05-17 14:41:55+00
hub	mission_portal_apache_from_stage	files	/var/cfengine/httpd/conf/httpd.conf	null	KEPT		2016-05-17 14:07:38+00
hub	cfe_internal_bins	files	/usr/local/sbin/cf-promises.cfsaved	null	KEPT		2016-05-17 13:13:33+00
host001	broken_promises	commands	/bin/false	demo_command_promise_not_kept	NOTKEPT	Executing 'no timeout' ... '/bin/false' Finished command related to promiser '/bin/false' -- an error occurred, returned 1 Completed execution of '/bin/false'	2016-05-17 14:41:55+00

Debugging

Inevitably, something will go wrong, and you will need to dig deep to figure something out. Lucky for you, I have some tips for debugging.

Run without locks

Again, using -K to disable locks is useful.

Using Verbose Mode

CFEngine's verbose output can be fantastic for debugging. Use the `-v` flag to turn it on.

```
cf-agent -Kv | grep -A 5 "BEGIN bundle"
```

When viewing verbose output, look for **BUNDLE <name>** for the bundle that you suspect is having trouble.

```
verbose: B: BEGIN bundle main
verbose: B: *****
verbose: P: .....
verbose: P: BEGIN promise 'promise_promises_cf_4' of type "reports" (pass 1)
verbose: P:   Promiser/affected object: 'Hello World!'
verbose: P:   Part of bundle: main
```

CFEngine will tell you exactly what is going on with each promise, in excruciating detail.

```
verbose: Using literal pathtype for '/tmp/touch'
verbose: No mode was set, choose plain file default 0600
  info: Created file '/tmp/touch', mode 0600
verbose: Handling file existence constraints on '/tmp/touch'
verbose: A: Promise REPAIRED
verbose: P: END files promise (/tmp/touch...)
```

Comments

CFEngine supports *comments* as part of its data structure. Every promise can have a comment attribute whose value is a quoted text string.

```
bundle agent example {  
  files:  
    "/etc/bind/named.cache"  
    copy_from => scp("${def.files}/bind/named.cache"),  
    comment   => "More recent copy of named.cache than shipped with bind";  
}
```

Comments show up in the verbose output.

```
verbose: P:      Container path : '/default/main/files'/etc/bind/named.cache'[0]'  
verbose: P:  
verbose: P:      Comment:  More recent copy of named.cache than shipped with bind.  
verbose: P: .....
```

The comment should always be **why** the promise is being made. Up until now none of the examples have used comments to save space on the slide. When writing your policies for real **every** promise should have a meaningful comment.

You'll thank me when this saves the day.

Promise Handles

When debugging, promise *handles* are also useful. Again, every promise can have a `handle` attribute whose value is a quoted canonical string.

```
bundle agent example{
  files:
    "/etc/bind/named.cache"
    copy_from => scp("${def.files}/bind/named.cache"),
    handle    => "update_etc_bind_named_cache",
    comment   => "More recent copy of named.cache than shipped with bind";
}
```

CFEngine will tell you the handle of each promise in the verbose output.

```
verbose: P: BEGIN promise 'update_etc_bind_named_cache' of type "files" (pass 1)
verbose: P:   Promiser/affected object: '/etc/bind/named.cache'
verbose: P:   Part of bundle: example
verbose: P:   Base context class: any
```

By giving each promise a unique handle you can swiftly jump back and forth between your debug output and your policy file. When writing your policies for real **every** promise should have a unique handle.

You'll thank me when this saves the day.

Promisees

When debugging, promise *stakeholders* aka *promisees* are useful for understanding who cares about a given promise.

```
bundle agent example {  
  files:  
    "/etc/bind/named.cache" -> { "Operations", "Nick Anderson" }  
    copy_from => scp("${def.files}/bind/named.cache"),  
    handle    => "update_etc_bind_named_cache",  
    comment   => "More recent copy of named.cache than shipped with bind";  
}
```

CFEngine will tell you additional info about each promise.

```
verbose: Additional promise info: handle 'update_etc_bind_named_cache'\  
       source path './t.cf' at line 4 promisee {'Operations','Nick Anderson'}\  
       comment 'More recent copy of named.cache than shipped with bind.'
```

Meta

When debugging variables and classes promise *meta* data is useful to help identify variables and classes with specific attributes.

debugging_classes_and_vars_with_tags.cf

```
bundle agent main{
  classes:
    "my_class" expression => "any", meta => { "mytag" };
  vars:
    "my_var" string => "value", meta => { "mytag" };
    "my_vars" slist => variablesmatching(".*", "mytag" );
    "my_classes" slist => classesmatching(".*", "mytag" );
  reports:
    "My var: $(my_vars)";
    "My class: $(my_classes)";
}
```

Note: Promise meta data is not currently displayed in the CFEngine's verbose output.

The Rest

Here's a list of topics that I didn't cover. This is to give you a taste of the rest of the power that is behind CFEngine. Dig deeper by checking them out in the [reference manual][reference-manual].

- `vars`: promises — Variables, strings, integers and reals (and lists of each).
- `methods`: promises — Create a self-contained bundle that can be called like a function.
- `guest_environments`: promises — Promise the existence of virtual machines.
- `storage`: promises — For local or remote (NFS) filesystems.
- `database` promises — Promise the schema of your database, CFEngine does the SQL for you.
- `edit_xml`: promises - Promise by path, CFEngine does the XML for you.
- Monitoring — Using data from `cf-monitor`.
- Implicit looping — Pass a list to a promise and it loops over the values in the list.

Pro Tips

- *Don't edit the standard library.* Create a `site_lib.cf` and add your custom library bundles and bodies there. This helps with upgrading because you won't have to patch your changes into the new version of the library. When you feel a bundle or body is ready for public use you can submit it to CFEngine by opening a pull request on [Github][masterfiles].
- *Make built-in classes and user defined classes easy to distinguish by sight.* CFEngine creates hard classes `all_lower_case_separated_by_underscores`. Whenever I define classes myself I use CamelCase.
- *Not sure how to organize masterfiles?* Check [A Case Study in CFEngine Layout][layout] by Brian Bennett.
- *Use `git` to revision control masterfiles.*
- *Syntax errors?* Only read the very first error. Fix it, then try again. A missing character in one promise will throw the whole file off.

[cfengine]: <http://github.com/cfengine/core> [masterfiles]:

<http://github.com/cfengine/masterfiles> [layout]:

<https://digitalelf.net/2013/04/a-case-study-in-cfengine-layout/> [reference-manual]:

<https://docs.cfengine.com/lts/reference.html>

Thanks

